

Kernel Modules

Kernel Module?

- Portion of kernel that can be dynamically loaded and unloaded
- Examples
 - USB drivers
 - File system drivers
 - Disk drivers
 - Cryptographic libraries

Why not just compile everything directly into the kernel?

- Each machine only needs a certain number of drivers
 - e.g. you don't need every single motherboard driver
- Load only the components you need
 - Smaller system footprint
 - Quicker boot time
- Dynamically load modules for new devices
 - New USB, camera, printer
 - Changing graphics card, motherboard, file system

Kernel Logistics

- Where to put kernel source tree
 - `/usr/src/<kernel name>`
 - Name yours something like `test_kernel` to make it easier to find
 - Issue your make commands here
- Where does the kernel image get installed to
 - `/boot/vmlinuz-<kernel name>`
 - Installed name might revert to kernel version (4.14.12 in this case)
- Where should I develop my kernel modules
 - `/usr/src/test_kernel/<module name>`
 - You can use symbolic links and point somewhere else to make things easier to find

Notes on Kernel Programming

- Kernel modules are event-driven
 - Register functions
 - Wait for requests from user-space and service them
 - Server/client model
- No standard C library
- No floating point support
- Crashes and deadlocks in a module could lead to crashing the entire kernel
 - Requires system-wide reboot

Creating a Kernel Module

hello.c

```
#include <linux/init.h>
#include <linux/module.h>
MODULE_LICENSE("Dual BSD/GPL");
static int hello_init(void) {
    printk(KERN_ALERT "Hello, world!\n");
    return 0;
}
static void hello_exit(void) {
    printk(KERN_ALERT "Goodbye, sleepy world.\n");
}
module_init(hello_init);
module_exit(hello_exit);
```

Creating a Kernel Module

hello.c

```
#include <linux/init.h>
#include <linux/module.h>
MODULE_LICENSE("Dual BSD/GPL");
static int hello_init(void) {
    printk(KERN_ALERT "Hello, world!\n");
    return 0;
}
static void hello_exit(void) {
    printk(KERN_ALERT "Goodbye, sleepy world.\n");
}
module_init(hello_init);
module_exit(hello_exit);
```



Module Headers

Creating a Kernel Module

hello.c

```
#include <linux/init.h>
#include <linux/module.h>
MODULE_LICENSE("Dual BSD/GPL");
static int hello_init(void) {
    printk(KERN_ALERT "Hello, world!\n");
    return 0;
}
static void hello_exit(void) {
    printk(KERN_ALERT "Goodbye, sleepy world.\n");
}
module_init(hello_init);
module_exit(hello_exit);
```



License Declaration

Creating a Kernel Module hello.c

```
#include <linux/init.h>
```

```
#include <linux/module.h>
```

```
MODULE_LICENSE("Dual BSD/GPL");
```

```
static int hello_init(void) {
```

```
    printk(KERN_ALERT "Hello, world!\n");
```

```
    return 0;
```

```
}
```

```
static void hello_exit(void) {
```

```
    printk(KERN_ALERT "Goodbye, sleepy world.\n");
```

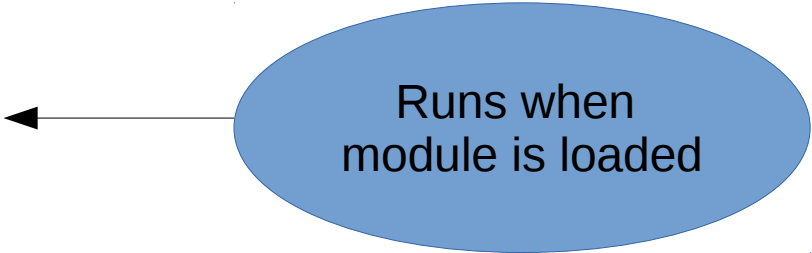
```
}
```

```
module_init(hello_init);
```

```
module_exit(hello_exit);
```



Initialization function



Runs when
module is loaded

Creating a Kernel Module hello.c

```
#include <linux/init.h>
#include <linux/module.h>
MODULE_LICENSE("Dual BSD/GPL");
static int hello_init(void) {
    printk(KERN_ALERT "Hello, world!\n");
    return 0;
}
static void hello_exit(void) {
    printk(KERN_ALERT "Goodbye, sleepy world.\n");
}
module_init(hello_init);
module_exit(hello_exit);
```

Exit function

A blue oval containing the text "Exit function" has an arrow pointing to the **module_exit(hello_exit);** line in the code block above.

Runs when
module is unloaded

A blue oval containing the text "Runs when module is unloaded" has an arrow pointing to the **module_exit(hello_exit);** line in the code block above.

Creating a Kernel Module Makefile

```
ifneq ($(KERNELRELEASE),)
    obj-m := hello.o
else
    KERNELDIR ?= /lib/modules/`uname -r`/build/
    PWD := `pwd`
default:
    $(MAKE) -C $(KERNELDIR) M=$(PWD) modules
endif
clean:
    rm -f *.ko *.o Module* *mod*
```

Creating a Kernel Module Compilation

```
/usr/src/test_kernel/hello $ make
```

Kernel Module Loading

- Insert a module

```
/usr/src/test_kernel/hello $ sudo insmod hello.ko
```

- Remove a module

```
/usr/src/test_kernel/hello $ sudo rmmod hello.ko
```

- List all running modules

```
/usr/src/test_kernel/hello $ lsmod
```

Kernel Functions

- `printf()` => `printk()`
- `malloc()` => `kmalloc()`
- `free()` => `kfree()`
- Where can I find definitions of these functions?
 - Section 9 of manpages
 - View online: <http://www.linuxsavvy.com/resources/linux/man/man9/>
 - Otherwise install from: <https://www.kernel.org/pub/linux/docs/man-pages/man-pages-4.14.tar.xz>

Kernel Headers

- `#include <linux/init.h>`
 - Module stuff
- `#include <linux/module.h>`
 - Module stuff
- `#include <asm/semaphore.h>`
 - Locks
- `#include <linux/list.h>`
 - Linked lists
- `#include <linux/string.h>`
 - String functions
- Look in `test_kernel/include/` for more
 - Search locally: `grep -Rn xtime /usr/src/test_kernel`
 - Search online: <http://lxr.free-electrons.com/>

printk

- Behaves similarly to printf
- Takes log level and format string as parameters
- Outputs to /var/log/syslog
 - \$ cat /var/log/syslog
 - \$ dmesg
- To watch syslog in realtime, use a second terminal to issue
 - \$ tail -f /var/log/syslog
- Can be called from just about anywhere at any time...
 - Except during booting before the console gets initialized

printk

- Log levels
 - KERN_EMERG Emergency condition, kernel likely crashed
 - KERN_ALERT Alert that requires immediate attention
 - KERN_CRIT Critical error message
 - KERN_ERR Error message
 - KERN_WARNING Warning message
 - KERN_NOTICE Normal, but noteworthy message
 - KERN_INFO Informational message
 - KERN_DEBUG Debug message
- Example
 - `printk(KERN_DEBUG "this is a debug message\n");`
 - Note there is no comma between the log level and the message