

Procfs Kernel Module

Procfs Kernel Module

- Hello World for /proc
- Steps
 - Create entry in module_init
 - Create file system functions
 - Open
 - Read
 - Close
 - Delete entry in module_cleanup

Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```

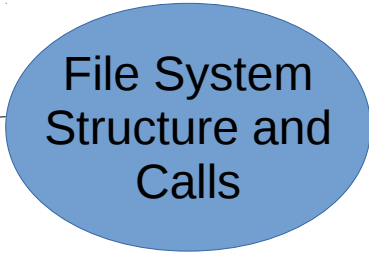
```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```

```
#define ENTRY_NAME "helloworld"
#define PERMS 0644
#define PARENT NULL
static struct file_operations fops;

static char *message;
static int read_p;
```

Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```



File System
Structure and
Calls

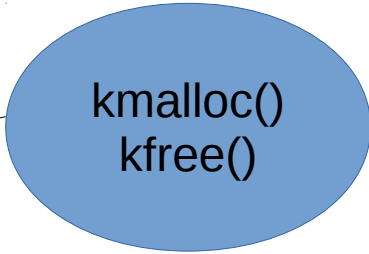
```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```

```
#define ENTRY_NAME "helloworld"
#define PERMS 0644
#define PARENT NULL
static struct file_operations fops;

static char *message;
static int read_p;
```

Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```



kmalloc()
kfree()

```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```

```
#define ENTRY_NAME "helloworld"
#define PERMS 0644
#define PARENT NULL
static struct file_operations fops;

static char *message;
static int read_p;
```

Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```



String
functions

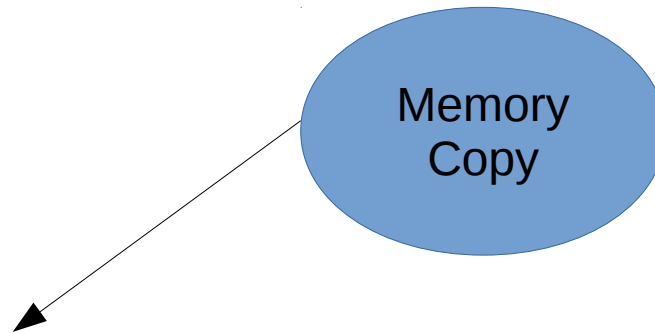
```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```

```
#define ENTRY_NAME "helloworld"
#define PERMS 0644
#define PARENT NULL
static struct file_operations fops;

static char *message;
static int read_p;
```

Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```



```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```

```
#define ENTRY_NAME "helloworld"
#define PERMS 0644
#define PARENT NULL
static struct file_operations fops;

static char *message;
static int read_p;
```

Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```

```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```



Module
Descriptions

```
#define ENTRY_NAME "helloworld"
#define PERMS 0644
#define PARENT NULL
static struct file_operations fops;

static char *message;
static int read_p;
```


Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```

```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```

```
#define ENTRY_NAME "helloworld"
```

```
#define PERMS 0644
```

```
#define PARENT NULL
```

```
static struct file_operations fops;
```

```
static char *message;
```

```
static int read_p;
```



Proc
Name

Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```

```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```

```
#define ENTRY_NAME "helloworld"
```

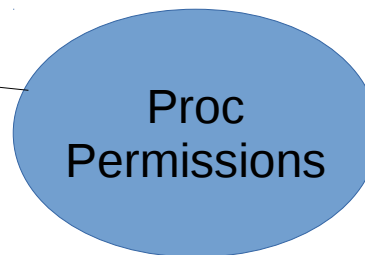
```
#define PERMS 0644
```

```
#define PARENT NULL
```

```
static struct file_operations fops;
```

```
static char *message;
```

```
static int read_p;
```



Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```

```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```

```
#define ENTRY_NAME "helloworld"
#define PERMS 0644
#define PARENT NULL
static struct file_operations fops;
```



Proc
Parent Directory

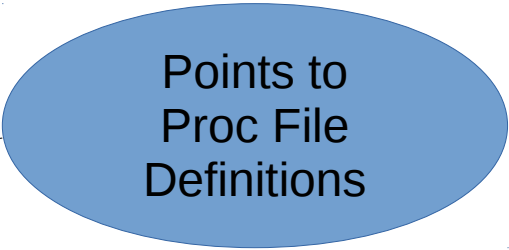
```
static char *message;
static int read_p;
```

Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```

```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```

```
#define ENTRY_NAME "helloworld"
#define PERMS 0644
#define PARENT NULL
static struct file_operations fops;
```



Points to
Proc File
Definitions

```
static char *message;
static int read_p;
```


Headers and Globals

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/proc_fs.h>
#include <linux/slab.h>
#include <linux/string.h>
#include <linux/uaccess.h>
```

```
MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple module featuring proc read");
```

```
#define ENTRY_NAME "helloworld"
#define PERMS 0644
#define PARENT NULL
static struct file_operations fops;
```

```
static char *message;
static int read_p;
```




Message to
Display in
Proc

Creation

```
int hello_proc_open(struct inode *sp_inode, struct file *sp_file) {  
    printk(KERN_INFO "proc called open\n");  
  
    read_p = 1;  
    message = kmalloc(sizeof(char) * 20, __GFP_RECLAIM |  
__GFP_WAIT | __GFP_IO | __GFP_FS);  
    if (message == NULL) {  
        printk(KERN_WARNING "hello_proc_open");  
        return -ENOMEM;  
    }  
    strcpy(message, "Hello, World!\n");  
    return 0;  
}
```

Creation

```
int hello_proc_open(struct inode *sp_inode, struct file *sp_file) {  
    printk(KERN_INFO "proc called open\n");  
  
    read_p = 1;  
    message = kmalloc(sizeof(char) * 20, __GFP_RECLAIM |  
__GFP_IO | __GFP_FS);  
    if (message == NULL) {  
        printk(KERN_WARNING "hello_proc_open");  
        return -ENOMEM;  
    }  
    strcpy(message, "Hello, World!\n");  
    return 0;  
}
```



Setup Proc Data Here

kmalloc()

- Takes
 - Number of bytes to allocate
 - A flag on how to allocate it
- Remember to restrict kernel memory allocation
 - Can block important functions
 - Can crash kernel if improperly handled
 - Kernel has limited access to memory

kmalloc()

- Flags

| | |
|---------------|--|
| __GFP_RECLAIM | Allocator can sleep |
| __GFP_HIGH | Allocator can access emergency pools |
| __GFP_IO | Allocator can start disk I/O |
| __GFP_FS | Allocator can start filesystem I/O |
| __GFP_COLD | Allocator should use cache cold pages |
| __GFP_NOWARN | Allocator will not print failure warnings |
| __GFP_REPEAT | Allocator will repeat if it fails (can still fail) |
| __GFP_NOFAIL | Allocator will repeat if it fails (can not fail) |
| __GFP_NORETRY | Allocator will never retry if it fails |
| __GFP_NO_GROW | Used by the slab |
| __GFP_COMP | Used by hugetlb |

Read

```
ssize_t hello_proc_read(struct file *sp_file, char __user *buf,  
size_t size, loff_t *offset) {  
    int len = strlen(message);  
  
    read_p = !read_p;  
    if (read_p)  
        return 0;  
  
    printk(KERN_INFO "proc called read\n");  
    copy_to_user(buf, message, len);  
    return len;  
}
```

Read

```
ssize_t hello_proc_read(struct file *sp_file, char __user *buf,  
size_t size, loff_t *offset) {
```

```
    int len = strlen(message);
```

```
    read_p = !read_p;
```

```
    if (read_p)
```

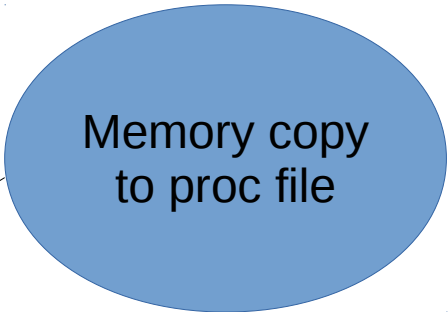
```
        return 0;
```

```
    printk(KERN_INFO "proc called read\n");
```

```
    copy_to_user(buf, message, len);
```

```
    return len;
```

```
}
```



Memory copy
to proc file

Read

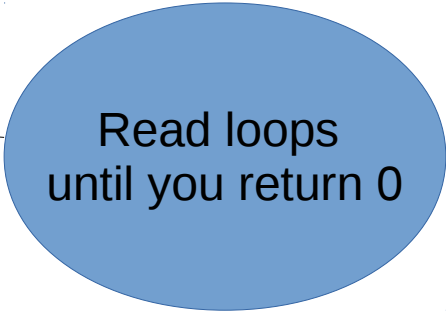
```
ssize_t hello_proc_read(struct file *sp_file, char __user *buf,  
size_t size, loff_t *offset) {
```

```
    int len = strlen(message);
```

```
    read_p = !read_p;
```

```
    if (read_p)
```

```
        return 0;
```



Read loops
until you return 0

```
    printk(KERN_INFO "proc called read\n");
```

```
    copy_to_user(buf, message, len);
```

```
    return len;
```

```
}
```

Memory Copying

Kernel → User

```
unsigned long copy_to_user (  
    void __user *to, const void *from, unsigned long size)
```

User → Kernel

```
unsigned long copy_from_user (  
    void *to, const void __user* from, unsigned long size)
```

- Needed because
 - User process uses virtual memory
 - Prevents crashing due to inaccessible regions
 - Can handle architecture specific issues

Close

```
int hello_proc_release(struct inode *sp_inode,  
struct file *sp_file) {  
    printk(KERN_INFO "proc called  
release\n");  
    kfree(message);  
    return 0;  
}
```

Init

```
static int hello_init(void) {
    printk(KERN_NOTICE "/proc/%s create\n", ENTRY_NAME);
    fops.open = hello_proc_open;
    fops.read = hello_proc_read;
    fops.release = hello_proc_release;

    if (!proc_create(ENTRY_NAME, PERMS, NULL, &fops)) {
        printk("ERROR! proc_create\n");
        remove_proc_entry(ENTRY_NAME, NULL);
        return -ENOMEM;
    }
    return 0;
}

module_init(hello_proc_init);
```

Init

```
static int hello_init(void) {  
    printk(KERN_NOTICE "/proc/%s create\n", ENTRY_NAME);  
    fops.open = hello_proc_open;  
    fops.read = hello_proc_read;  
    fops.release = hello_proc_release;  
  
    if (!proc_create(ENTRY_NAME, PERMS, NULL, &fops)) {  
        printk("ERROR! proc_create\n");  
        remove_proc_entry(ENTRY_NAME, NULL);  
        return -ENOMEM;  
    }  
    return 0;  
}  
  
module_init(hello_proc_init);
```



Setup
Proc calls

Init

```
static int hello_init(void) {  
    printk(KERN_NOTICE "/proc/%s create\n", ENTRY_NAME);  
    fops.open = hello_proc_open;  
    fops.read = hello_proc_read;  
    fops.release = hello_proc_release;  
  
    if (!proc_create(ENTRY_NAME, PERMS, NULL, &fops)) {  
        printk("ERROR! proc_create\n");  
        remove_proc_entry(ENTRY_NAME, NULL);  
        return -ENOMEM;  
    }  
    return 0;  
}  
module_init(hello_proc_init);
```



Make
Proc Entry

Exit


```
static void hello_exit(void) {  
    remove_proc_entry(ENTRY_NAME,  
NULL);  
    printk(KERN_NOTICE "Removing /proc/  
%s.\n", ENTRY_NAME);  
}
```

```
module_exit(hello_proc_exit);
```

Exit

```
static void hello_exit(void) {  
    remove_proc_entry(ENTRY_NAME,  
NULL);  
    printk(KERN_NOTICE "Removing /proc/  
%s.\n", ENTRY_NAME);  
}
```

```
module_exit(hello_proc_exit);
```



Remove
Proc entry

Testing

```
$ make
```

```
$ sudo insmod hello_proc.ko
```

```
$ dmesg | tail
```

```
$ cat /proc/helloworld
```

```
$ sudo rmmod hello_proc
```

```
$ dmesg | tail
```