

# Scheduling

# Why Make an Elevator?

- It demonstrates the producer / consumer analogy
  - People are produced
    - Through system calls
  - The elevator consumes them
    - By taking them to the appropriate floor

# Examples

- File system
  - Users produce read/write requests
  - Disk consumes the requests by returning the file data
    - Prioritizes based on disk head position and rotational delay (throughput)
- Web
  - Clients request to view web pages
  - Server consumes requests by returning the page data
    - Prioritizes based on decreasing the amount of time each client has to wait (latency)

# Simple Design

- While elevator is on
  - Pickup as many people as you can on a floor
  - Loop while not empty
    - Go to the destination of the first person
    - Drop off as many people as you can
      - Optionally, pick people up from that floor as well
- Start with something like this
- Optimize time permitting

# Scheduler Design Choices

- Different tasks have different needs
  - There is no one size-fits-all model
- Some metrics that can be optimized (not mutually exclusive)
  - Throughput
    - Maximize total number of requests serviced
  - Burst throughput
    - Be able to handle bursts of requests
  - Latency
    - Minimize time it takes to service each request
  - Priorities
    - Service requests with higher priorities first
  - Energy
    - Minimize processing time
    - Different than the time to service the requests

# Scheduler Design Choices

- Elevator scheduler should be optimized for throughput
  - You don't need to worry about
    - How long it takes for each person
    - Priorities among different types of people
    - Priorities for which floor to service
    - Amount of effort spent loading vs moving

# Scheduler Algorithms

## FIFO

- Method
  - Service requests in the order they arrive
- Pros
  - No computational time needed
  - Easy implementation
  - Every request gets a turn
- Cons
  - Bad at random requests
    - Elevator needs to go back and forth several times in processing requests
  - Low throughput potential

# Scheduler Algorithms

## Shortest Seek Time First

- Method
  - Service requests near elevator first
- Pros
  - Very low computational time expended
  - Implementation is fairly easy
  - Fairly high throughput by minimizing seek time
- Cons
  - Unfair to distant requests
    - Note, you're trying to maximize throughput, you're not worrying about the individual requests
  - May ignore distant clusters of requests



# Scheduler Algorithms

## SCAN

- Method
  - Service requests in one direction, then go backwards
- Pros
  - Very low computational time expended
  - Implementation is easy (up, down, up, down, ...)
  - Throughput is somewhat high
  - Fair due to no starvation
- Cons
  - May take up to two trips to collect a set of requests (up,down)

# Scheduler Algorithms

## LOOK

- Method
  - Improvement on SCAN
  - Uses information of request locations to determine where to stop going in a certain direction
- Pros
  - Low computational time expended
  - Moderate implementation difficulty
    - Set upper bound, go up, set lower bound, go down
  - Throughput is somewhat high
  - Fair due to no starvation
- Cons
  - May miss a new request at outer boundary just as direction change occurs

# Scheduler Algorithms

## Hybrid

- Combine methods and come up with something new
- Up to your creativity
- Example
  - Shortest Seek Time First w/ Scan
  - Do shortest seek time on a window size of  $k$
  - Move the window up or down over after  $j$  ticks
  - e.g. SSTF for floors 1-3, then 2-4, then 3-6, then 2-4, ... (for window of 3, 6 floors total)